

dog_app

June 16, 2020

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

Note: if you are using the Udacity workspace, you DO NOT need to re-download these - they can be found in the /data folder as noted in the cell below.

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human_files and dog_files.

```
In [1]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        human_files = np.array(glob("/data/lfw/*/.*"))
        dog_files = np.array(glob("/data/dog_images/*/.*"))

        # print number of images in each dataset
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the haarcascades directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [2]: import cv2
        import matplotlib.pyplot as plt
        %matplotlib inline

        # extract pre-trained face detector
        face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

        # load color (BGR) image
        img = cv2.imread(human_files[0])
        # convert BGR image to grayscale
        gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

        # find faces in image
        faces = face_cascade.detectMultiScale(gray)

        # print number of faces detected in the image
        print('Number of faces detected:', len(faces))
```

```

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [3]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer: (You can print out your results and/or write your percentages in this cell)

```
In [4]: from tqdm import tqdm

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

#-#-# Do NOT modify the code above this line. #-#-#

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.

human_files_as_human = np.average([face_detector(img) for img in tqdm(human_files_short)])
dog_files_as_human = np.average([face_detector(img) for img in tqdm(dog_files_short)])

print('Accuracy human Avg : {}'.format(human_files_as_human))
print('Dog Avg is detected as human : {}'.format(dog_files_as_human))
```

```
100%|| 100/100 [00:02<00:00, 35.95it/s]
100%|| 100/100 [00:29<00:00, 7.34it/s]
```

Accuracy human Avg : 0.98

Dog Avg is detected as human : 0.17

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [25]: ### (Optional)  
        ### TODO: Test performance of another face detection algorithm.  
        ### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```
In [5]: import torch  
        import torchvision.models as models  
  
        # define VGG16 model  
        VGG16 = models.vgg16(pretrained=True)  
  
        # check if CUDA is available  
        use_cuda = torch.cuda.is_available()  
  
        # move model to GPU if CUDA is available  
        if use_cuda:  
            VGG16 = VGG16.cuda()
```

```
Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg16-397923af.pth  
100%|| 553433881/553433881 [00:05<00:00, 100650407.03it/s]
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as `'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg'`) as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```

In [6]: from PIL import Image
import torchvision.transforms as transforms

def VGG16_predict(img_path):
    '''
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    '''

    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    ## Return the *index* of the predicted class for that image

    # Open the image
    img = Image.open(img_path)

    # convert image into Tensor using (toTensor = transforms.ToTensor())

    # human face width=250 (jpg file), dog jpg file size is different>> resize into 250

    transform_pipeline = transforms.Compose([transforms.RandomResizedCrop(250),
                                             transforms.ToTensor()])

    img_tensor = transform_pipeline(img)
    img_tensor = img_tensor.unsqueeze(0)

    # from tensor to cuda
    if torch.cuda.is_available():
        img_tensor = img_tensor.cuda()

    prediction = VGG16(img_tensor)

    # from tensor to cpu (for cpu processing)
    if torch.cuda.is_available():
        prediction = prediction.cpu()

    index = prediction.data.numpy().argmax()

    return index # predicted class index

```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns True if a dog is detected in an image (and False if not).

```
In [7]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    ## TODO: Complete the function.
    index = VGG16_predict(img_path)

    # print index between 151 and 268 (inclusive)
    return (151 <= index & index <= 268) # true/false
```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

Answer:

```
In [8]: ### TODO: Test the performance of the dog_detector function
        ### on the images in human_files_short and dog_files_short.

human_files_as_human = np.average([dog_detector(img) for img in human_files_short])
dog_files_as_human = np.average([dog_detector(img) for img in dog_files_short])

print('Human Avg is detected as Dog : {}'.format(human_files_as_human))
print('Accuracy Dog Avg : {}'.format(dog_files_as_human))
```

Human Avg is detected as Dog : 0.01

Accuracy Dog Avg : 0.77

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [30]: ### (Optional)
        ### TODO: Report the performance of another pre-trained network.
        ### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [9]: import os
        from torchvision import datasets
        from PIL import ImageFile
        import torchvision.transforms as transforms

        ImageFile.LOAD_TRUNCATED_IMAGES = True
```



```
### TODO: Write data loaders for training, validation, and test sets
## Specify appropriate transforms, and batch_sizes
```

```
# Pre-process train dataset with augmentation
```

```
train_transform = transforms.Compose([
    transforms.RandomResizedCrop(224),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225]))
```

```
# Pre-process valid/test dataset without augmentation
```

```
transform_pipeline = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225]))
```

```
train_data = datasets.ImageFolder('/data/dog_images/train/', transform=train_transform)
valid_data = datasets.ImageFolder('/data/dog_images/valid/', transform=transform_pipeline)
test_data = datasets.ImageFolder('/data/dog_images/test/', transform=transform_pipeline)
```

```
batch_size = 10
num_workers = 0
```

```
train_loader = torch.utils.data.DataLoader(train_data,
                                             batch_size=batch_size,
                                             num_workers=num_workers,
                                             shuffle=True)
```

```
valid_loader = torch.utils.data.DataLoader(valid_data,
                                             batch_size=batch_size,
                                             num_workers=num_workers,
                                             shuffle=False)
```

```
test_loader = torch.utils.data.DataLoader(test_data,
                                             batch_size=batch_size,
                                             num_workers=num_workers,
                                             shuffle=False)
```

```
loaders_scratch = {
    'train': train_loader,
    'valid': valid_loader,
```

```

        'test': test_loader
    }

```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer:

- **How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why?**

I have use (RandomResizedCrop) with scaled size into 224

- **Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?**

Yes, by Convert image file to tensor

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [10]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture

dog_classes = 133 # total dog class

class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN
        self.conv1 = nn.Conv2d(3, 32, 3, padding=1)
        self.norm2d1 = nn.BatchNorm2d(32)
        self.conv2 = nn.Conv2d(32, 64, 3, padding=1)
        self.conv3 = nn.Conv2d(64, 128, 3, padding=1)

        # pooling
        self.pool = nn.MaxPool2d(2, 2)

        size_linear_layer = 500

        # linear layer (128 * 28 * 28 ==> 500)
        self.fc1 = nn.Linear(128 * 28 * 28, size_linear_layer)
        self.fc2 = nn.Linear(size_linear_layer, dog_classes)

```

```

def forward(self, x):
    ## Define forward behavior
    x = self.pool(F.relu(self.norm2d1(self.conv1(x))))
    x = self.pool(F.relu(self.conv2(x)))
    x = self.pool(F.relu(self.conv3(x)))

    # image input
    x = x.view(-1, 128 * 28 * 28)
    x = F.relu(self.fc1(x))
    x = self.fc2(x)
    return x

### You so NOT have to modify the code below this line. ###

# instantiate the CNN
model_scratch = Net()
print(model_scratch)

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch = model_scratch.cuda()

Net(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (norm2d1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): Linear(in_features=100352, out_features=500, bias=True)
  (fc2): Linear(in_features=500, out_features=133, bias=True)
)

```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer:

1. Conv 1: layer, activation, pooling

-(conv1): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

-Activation Fun: relu

-(pooling): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

2. Conv 2: layer, activation, pooling

-(conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

-Activation Fun:relu

-(pooling): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

3. Conv 3: layer, activation, pooling

- (conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
- Activation Fun:relu
- (pooling): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

4. (fc1): Linear(in_features=61504, out_features=500, bias=True)

5. (fc2): Linear(in_features=500, out_features=133, bias=True)

There are three conv layers, and maxpooling

I have add (Maxpool) when image passed cov layer to avoid iverfitting in each cov layer and to decrease the featrue mape.

in each cov layer, width and height are decreased as below:

(250,250) = input image size

-> (125,125)

-> (62,62)

-> (28,28)

The depth = 64 ==> i.e: the fully-connected layer size is $64 * 28 * 28$.

Also, It is connected to last fully connected layer and the size=133 which is similar to the total classes of dog

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```
In [11]: import torch.optim as optim

        ### TODO: select loss function
        criterion_scratch = nn.CrossEntropyLoss()

        ### TODO: select optimizer
        optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=0.01)

        if use_cuda:
            criterion_scratch.cuda()
```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```
In [12]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
        """returns trained model"""
        # initialize tracker for minimum validation loss
        valid_loss_min = np.Inf

        for epoch in range(1, n_epochs+1):
```

```

# initialize variables to monitor training and validation loss
train_loss = 0.0
valid_loss = 0.0

#####
# train the model #
#####
model.train()
for batch_idx, (data, target) in enumerate(loaders['train']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    ## find the loss and update the model parameters accordingly
    ## record the average training loss, using something like
    ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

    optimizer.zero_grad()

    # forward pass
    output = model(data)

    # calculate batch loss
    loss = criterion(output, target)

    # backward pass
    loss.backward()

    # parameter update
    optimizer.step()

    train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

#####
# validate the model #
#####
model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    ## update the average validation loss
    output = model(data)
    loss = criterion(output, target)
    valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

# print training/validation statistics

```

```

print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

## TODO: save the model if validation loss has decreased
if valid_loss < valid_loss_min:
    torch.save(model.state_dict(), save_path)
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'
          .format(valid_loss_min, valid_loss))
    valid_loss_min = valid_loss

# return trained model
return model

# train the model
model_scratch = train(35, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, 'model_scratch.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))

```

```

Epoch: 1      Training Loss: 4.874071      Validation Loss: 4.799191
Validation loss decreased (inf --> 4.799191). Saving model ...
Epoch: 2      Training Loss: 4.690374      Validation Loss: 4.483826
Validation loss decreased (4.799191 --> 4.483826). Saving model ...
Epoch: 3      Training Loss: 4.579238      Validation Loss: 4.438984
Validation loss decreased (4.483826 --> 4.438984). Saving model ...
Epoch: 4      Training Loss: 4.518841      Validation Loss: 4.358006
Validation loss decreased (4.438984 --> 4.358006). Saving model ...
Epoch: 5      Training Loss: 4.451295      Validation Loss: 4.277254
Validation loss decreased (4.358006 --> 4.277254). Saving model ...
Epoch: 6      Training Loss: 4.382032      Validation Loss: 4.201300
Validation loss decreased (4.277254 --> 4.201300). Saving model ...
Epoch: 7      Training Loss: 4.324695      Validation Loss: 4.123540
Validation loss decreased (4.201300 --> 4.123540). Saving model ...
Epoch: 8      Training Loss: 4.287301      Validation Loss: 4.111712
Validation loss decreased (4.123540 --> 4.111712). Saving model ...
Epoch: 9      Training Loss: 4.219664      Validation Loss: 4.080149
Validation loss decreased (4.111712 --> 4.080149). Saving model ...
Epoch: 10     Training Loss: 4.164010      Validation Loss: 4.054220
Validation loss decreased (4.080149 --> 4.054220). Saving model ...
Epoch: 11     Training Loss: 4.142148      Validation Loss: 3.987947
Validation loss decreased (4.054220 --> 3.987947). Saving model ...
Epoch: 12     Training Loss: 4.076387      Validation Loss: 3.985136
Validation loss decreased (3.987947 --> 3.985136). Saving model ...

```

```

Epoch: 13      Training Loss: 4.042631      Validation Loss: 3.866470
Validation loss decreased (3.985136 --> 3.866470). Saving model ...
Epoch: 14      Training Loss: 3.996782      Validation Loss: 3.832814
Validation loss decreased (3.866470 --> 3.832814). Saving model ...
Epoch: 15      Training Loss: 3.936854      Validation Loss: 3.856555
Epoch: 16      Training Loss: 3.895902      Validation Loss: 3.760493
Validation loss decreased (3.832814 --> 3.760493). Saving model ...
Epoch: 17      Training Loss: 3.855852      Validation Loss: 3.739038
Validation loss decreased (3.760493 --> 3.739038). Saving model ...
Epoch: 18      Training Loss: 3.793676      Validation Loss: 3.703358
Validation loss decreased (3.739038 --> 3.703358). Saving model ...
Epoch: 19      Training Loss: 3.752026      Validation Loss: 3.742903
Epoch: 20      Training Loss: 3.721624      Validation Loss: 3.705671
Epoch: 21      Training Loss: 3.673020      Validation Loss: 3.700200
Validation loss decreased (3.703358 --> 3.700200). Saving model ...
Epoch: 22      Training Loss: 3.646251      Validation Loss: 3.659706
Validation loss decreased (3.700200 --> 3.659706). Saving model ...
Epoch: 23      Training Loss: 3.586033      Validation Loss: 3.621545
Validation loss decreased (3.659706 --> 3.621545). Saving model ...
Epoch: 24      Training Loss: 3.528124      Validation Loss: 3.607066
Validation loss decreased (3.621545 --> 3.607066). Saving model ...
Epoch: 25      Training Loss: 3.510861      Validation Loss: 3.567182
Validation loss decreased (3.607066 --> 3.567182). Saving model ...
Epoch: 26      Training Loss: 3.444341      Validation Loss: 3.677048
Epoch: 27      Training Loss: 3.416143      Validation Loss: 3.571202
Epoch: 28      Training Loss: 3.351775      Validation Loss: 3.588123
Epoch: 29      Training Loss: 3.313012      Validation Loss: 3.512544
Validation loss decreased (3.567182 --> 3.512544). Saving model ...
Epoch: 30      Training Loss: 3.256693      Validation Loss: 3.561528
Epoch: 31      Training Loss: 3.239824      Validation Loss: 3.580557
Epoch: 32      Training Loss: 3.190357      Validation Loss: 3.514172
Epoch: 33      Training Loss: 3.140689      Validation Loss: 3.451639
Validation loss decreased (3.512544 --> 3.451639). Saving model ...
Epoch: 34      Training Loss: 3.098761      Validation Loss: 3.553817
Epoch: 35      Training Loss: 3.083761      Validation Loss: 3.499618

```

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```

In [13]: def test(loaders, model, criterion, use_cuda):

        # monitor test loss and accuracy
        test_loss = 0.
        correct = 0.
        total = 0.

```

```

model.eval()
for batch_idx, (data, target) in enumerate(loaders['test']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the loss
    loss = criterion(output, target)
    # update average test loss
    test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
    # convert output probabilities to predicted class
    pred = output.data.max(1, keepdim=True)[1]
    # compare predictions to true label
    correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
    total += data.size(0)

print('Test Loss: {:.6f}\n'.format(test_loss))

print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
    100. * correct / total, correct, total))

# call test function
test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)

```

Test Loss: 3.387951

Test Accuracy: 20% (171/836)

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```

In [14]: ## TODO: Specify data loaders
         loaders_transfer = loaders_scratch.copy()

```


1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```
In [15]: import torchvision.models as models
import torch.nn as nn

## TODO: Specify model architecture
model_transfer = models.resnet50(pretrained=True)

for param in model_transfer.parameters():
    param.requires_grad = False

model_transfer.fc = nn.Linear(2048, 133, bias=True)

fc_parameters = model_transfer.fc.parameters()

for param in fc_parameters:
    param.requires_grad = True

if use_cuda:
    model_transfer = model_transfer.cuda()
```

Downloading: "https://download.pytorch.org/models/resnet50-19c8e357.pth" to /root/.torch/models/100%|| 102502400/102502400 [00:01<00:00, 57652152.55it/s]

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer:

For good image classification (ResNet) was chosen.

At the end, fully-connected layer is added plus a fully-connected layer with output of 133 (which represent the total dog classes).

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [16]: criterion_transfer = nn.CrossEntropyLoss()
optimizer_transfer = optim.SGD(model_transfer.fc.parameters(), lr=0.001)
```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```

In [17]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
           """returns trained model"""

           # initialize tracker for minimum validation loss
           valid_loss_min = np.Inf

           for epoch in range(1, n_epochs+1):

               # initialize variables to monitor training and validation loss
               train_loss = 0.0
               valid_loss = 0.0

               #####
               # train the model #
               #####

               model.train()

               for batch_idx, (data, target) in enumerate(loaders['train']):
                   # move to GPU
                   if use_cuda:
                       data, target = data.cuda(), target.cuda()
                   ## find the loss and update the model parameters accordingly
                   ## record the average training loss, using something like
                   ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

                   optimizer.zero_grad()

                   # forward pass
                   output = model(data)

                   # calculate batch loss
                   loss = criterion(output, target)

                   # backward pass
                   loss.backward()

                   # parameter update
                   optimizer.step()

                   train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

               #####
               # validate the model #
               #####
               model.eval()
               for batch_idx, (data, target) in enumerate(loaders['valid']):

```

```

        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        ## update the average validation loss
        output = model(data)
        loss = criterion(output, target)
        valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

    # print training/validation statistics
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
        epoch,
        train_loss,
        valid_loss
    ))

    ## TODO: save the model if validation loss has decreased
    if valid_loss < valid_loss_min:
        torch.save(model.state_dict(), save_path)
        print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'
              .format(valid_loss_min, valid_loss))
        valid_loss_min = valid_loss

    # return trained model
    return model

n_epochs = 20

# train the model
# train(n_epochs, loaders_transfer, model_transfer, optimizer_transfer,
#       criterion_transfer, use_cuda, 'model_transfer.pt')
model_transfer = train(n_epochs, loaders_transfer, model_transfer, optimizer_transfer,
                       use_cuda, 'model_transfer.pt')

# load the model that got the best validation accuracy (uncomment the line below)
model_transfer.load_state_dict(torch.load('model_transfer.pt'))

```

```

Epoch: 1          Training Loss: 4.745024          Validation Loss: 4.443148
Validation loss decreased (inf --> 4.443148). Saving model ...
Epoch: 2          Training Loss: 4.429244          Validation Loss: 4.044027
Validation loss decreased (4.443148 --> 4.044027). Saving model ...
Epoch: 3          Training Loss: 4.148744          Validation Loss: 3.653620
Validation loss decreased (4.044027 --> 3.653620). Saving model ...
Epoch: 4          Training Loss: 3.881971          Validation Loss: 3.302384
Validation loss decreased (3.653620 --> 3.302384). Saving model ...
Epoch: 5          Training Loss: 3.641710          Validation Loss: 2.999373
Validation loss decreased (3.302384 --> 2.999373). Saving model ...

```

```

Epoch: 6      Training Loss: 3.422979      Validation Loss: 2.757210
Validation loss decreased (2.999373 --> 2.757210). Saving model ...
Epoch: 7      Training Loss: 3.236698      Validation Loss: 2.506831
Validation loss decreased (2.757210 --> 2.506831). Saving model ...
Epoch: 8      Training Loss: 3.040642      Validation Loss: 2.304289
Validation loss decreased (2.506831 --> 2.304289). Saving model ...
Epoch: 9      Training Loss: 2.893299      Validation Loss: 2.098266
Validation loss decreased (2.304289 --> 2.098266). Saving model ...
Epoch: 10     Training Loss: 2.757908      Validation Loss: 1.949198
Validation loss decreased (2.098266 --> 1.949198). Saving model ...
Epoch: 11     Training Loss: 2.626062      Validation Loss: 1.819421
Validation loss decreased (1.949198 --> 1.819421). Saving model ...
Epoch: 12     Training Loss: 2.517177      Validation Loss: 1.724380
Validation loss decreased (1.819421 --> 1.724380). Saving model ...
Epoch: 13     Training Loss: 2.420096      Validation Loss: 1.606547
Validation loss decreased (1.724380 --> 1.606547). Saving model ...
Epoch: 14     Training Loss: 2.320333      Validation Loss: 1.463492
Validation loss decreased (1.606547 --> 1.463492). Saving model ...
Epoch: 15     Training Loss: 2.243379      Validation Loss: 1.379712
Validation loss decreased (1.463492 --> 1.379712). Saving model ...
Epoch: 16     Training Loss: 2.146862      Validation Loss: 1.328910
Validation loss decreased (1.379712 --> 1.328910). Saving model ...
Epoch: 17     Training Loss: 2.077968      Validation Loss: 1.260870
Validation loss decreased (1.328910 --> 1.260870). Saving model ...
Epoch: 18     Training Loss: 2.020319      Validation Loss: 1.200652
Validation loss decreased (1.260870 --> 1.200652). Saving model ...
Epoch: 19     Training Loss: 1.956258      Validation Loss: 1.124372
Validation loss decreased (1.200652 --> 1.124372). Saving model ...
Epoch: 20     Training Loss: 1.932868      Validation Loss: 1.108843
Validation loss decreased (1.124372 --> 1.108843). Saving model ...

```

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [18]: test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

```
Test Loss: 1.084870
```

```
Test Accuracy: 80% (673/836)
```

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```

In [21]: ### TODO: Write a function that takes a path to an image as input
         ### and returns the dog breed that is predicted by the model.

from PIL import Image
import torchvision.transforms as transforms

data_transfer = loaders_transfer.copy()

# list of class names by index, i.e. a name can be accessed like class_names[0]
class_names = [item[4:].replace("_", " ") for item in data_transfer['train'].dataset.cl

def predict_breed_transfer(img_path):
    global model_transfer
    global train_transform

    # load the image and return the predicted breed
    image = Image.open(img_path).convert('RGB')

    # Removing transparent, alpha
    image = train_transform(image)[:3,:,:].unsqueeze(0)

    if use_cuda:
        model_transfer = model_transfer.cuda()
        image = image.cuda()

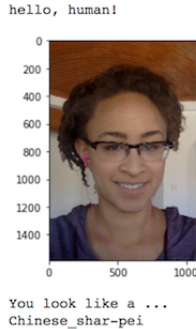
    model_transfer.eval()
    idx = torch.argmax(model_transfer(image))
    return class_names[idx]

In [22]: # For test predict_breed_transfer

for img_file in os.listdir("/data/dog_images/test/001.Affenpinscher/"):
    img_path = os.path.join('/data/dog_images/test/001.Affenpinscher/', img_file)
    predition = predict_breed_transfer(img_path)
    print("image_file_name: {0}, \t predition breed: {1}".format(img_path, predition))

image_file_name: /data/dog_images/test/001.Affenpinscher/Affenpinscher_00036.jpg,      predi
image_file_name: /data/dog_images/test/001.Affenpinscher/Affenpinscher_00047.jpg,      predi
image_file_name: /data/dog_images/test/001.Affenpinscher/Affenpinscher_00071.jpg,      predi
image_file_name: /data/dog_images/test/001.Affenpinscher/Affenpinscher_00023.jpg,      predi
image_file_name: /data/dog_images/test/001.Affenpinscher/Affenpinscher_00048.jpg,      predi
image_file_name: /data/dog_images/test/001.Affenpinscher/Affenpinscher_00058.jpg,      predi
image_file_name: /data/dog_images/test/001.Affenpinscher/Affenpinscher_00078.jpg,      predi

```



Sample Human Output

image_file_name: /data/dog_images/test/001.Affenpinscher/Affenpinscher_00003.jpg,

predi

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.1.18 (IMPLEMENTATION) Write your Algorithm

```
In [23]: ### TODO: Write your algorithm.
        ### Feel free to use as many code cells as needed.

def run_app(img_path):

    ## handle cases for a human face, dog, and neither
    if face_detector(img_path) > 0:
        breed = predict_breed_transfer(img_path)
        print('Human / similar to dog breed is ' + breed)
    elif dog_detector(img_path):
        breed = predict_breed_transfer(img_path)
        print('Dog / dog breed is ' + breed)
    else:
        print('Not Dog, Not Human')
```

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: (Three possible points for improvement)

- Adding more breeds of dogs and/or increase human pictures
- Increase number of epochs (if we have great resource will help to discover)
- Playing with some parameters could help too, such as inputs and output values (maybe will affect)

```
In [24]: import os
         from PIL import Image


         ## TODO: Execute your algorithm from Step 6 on
         ## at least 6 images on your computer.
         ## Feel free to use as many code cells as needed.

         for img_file in os.listdir('/home/workspace/dog_project/new_images/'):
             img_path = os.path.join('/home/workspace/dog_project/new_images/', img_file)
             run_app(img_path)
             print(img_path)
             img = Image.open(img_path)
             plt.imshow(img)
             plt.show()


Dog / dog breed is Akita
/home/workspace/dog_project/new_images/ThinkstockPhotos-471884456.jpg
```




Dog / dog breed is Afghan hound
/home/workspace/dog_project/new_images/2. Afghan Hound.jpg




Human / similar to dog breed is Italian greyhound
/home/workspace/dog_project/new_images/sub-buzz-1096-1579879662-3.jpg




Dog / dog breed is Alaskan malamute
/home/workspace/dog_project/new_images/0_EK8EubWpGkk72RIT.png



Human / similar to dog breed is Dogue de bordeaux
/home/workspace/dog_project/new_images/barack-obama-writers-photo-1.jpeg




Dog / dog breed is Chinese crested
/home/workspace/dog_project/new_images/real-names-of-famous-people-1.jpg




Human / similar to dog breed is Maltese

/home/workspace/dog_project/new_images/famous-people-with-anxiety-rm-selena-gomez-722x406.jpg



Dog / dog breed is Chow chow

/home/workspace/dog_project/new_images/f75335ae6328ae3556cb3e8c5046bfc9.jpg



In []: